

Modern Normalising Flows

Denis Derkach

CS HSE faculty, Generative Models, spring 2020

Contents

Basic Understanding

Masked Autoregressive Flows

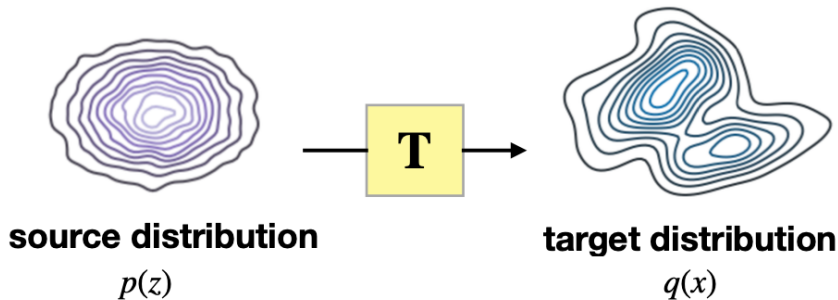
GLOW

FFJORD

Basic Understanding

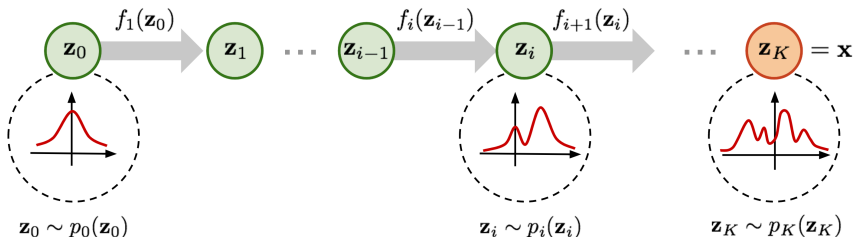
Reminder: Motivation

Quite often we want to sample from some distribution (but we do not know how to do this). What we know is how to sample from a simple pdf: Gaussian, uniform or something.



We want a deterministic map T from source to target density.

Definition



For definition:

$$z_m = f_\theta^m \circ f_\theta^{m-1} \dots \circ f_\theta^1(z_0) = f_\theta^m(f_\theta^{m-1}(\dots(f_\theta^1(z_0)))) = f_\theta(z_0),$$

we have $z_m \rightarrow x$:

$$p(\mathbf{x}; \theta) = q(f_\theta^{-1}(\mathbf{z})) \prod_{i=1}^m \left| \det \frac{\partial (f_\theta^i)^{-1}(\mathbf{z}^m)}{\partial \mathbf{z}^m} \right|.$$

Note that we have invertible transformations.

Triangular Jacobian

- › There always exists a unique (up to ordering) increasing triangular map that transforms a source density to a target density (see Bogachev et al. for detail).
- › In previous lectures we considered the triangular Jacobian:

$$J = \begin{pmatrix} \frac{\partial f_1}{\partial z_1} & \cdots & 0 \\ \cdots & \cdots & \cdots \\ \frac{\partial f_1}{\partial z_n} & \cdots & \frac{\partial f_n}{\partial z_n} \end{pmatrix}$$

It describes the flow with transformation $x_i = f_i(z)$ that only depends on $z \leq i$.

Developments

More results are produced in 2019. In general, they can be separated into following categories.

1. Det Identities

Planar NF
Sylvester NF
...

Jacobian



(Low rank)

2. Coupling Blocks

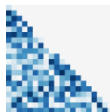
NICE
Real NVP
Glow
...



(Lower triangular +
structured)

3. Autoregressive

Inverse AF
Neural AF
Masked AF
...



(Lower triangular)

4. Unbiased Estimation

FFJORD
Residual Flows



(Arbitrary)

Masked Autoregressive Flows

Reminder: Autoregressive models

- › Take Autoregressive Model:

$$p(x) = \prod_{i=1}^n p(x_i | x_{i < 1}).$$

such that

$$p(x_i | x_{i < 1}) = \mathcal{N}(\mu_i(x_1, \dots, x_{i-1}), \exp(\alpha_i(x_1, \dots, x_{i-1}))^2),$$

with μ and α are Neural network outputs.

- › We have a direct estimation of likelihood in this model.
- › To sample, we need to go through consecutive steps:
 - › $z_i \sim \mathcal{N}(0; 1)$;
 - › $x_1 = \exp(\alpha_1)z_1 + \mu_1$
 - › $x_2 = \exp(\alpha_2(x_1))z_1 + \mu_2(x_1)$
 - › and so on.
- › Might be stacked as Flow from Gaussians to observable space.

Masked and Inverse Autoregressive Flow (MAF / IAF)

- › looks similar to MADE;
- › Forward mapping $z \mapsto x$:

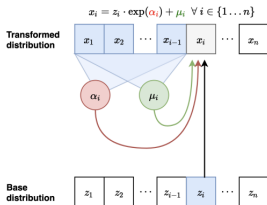
- › $z_i \sim \mathcal{N}(0; 1)$;

- › $x_1 = \exp(\alpha_1)z_1 + \mu_1$

- › $x_2 = \exp(\alpha_2(x_1))z_2 + \mu_2(x_1)$

- › and so on.

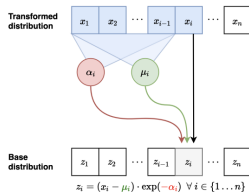
- › sampling is sequential and slow.



From Stanford GM lectures by S. Ermano G. Papamakarios et al. Masked Autoregressive Flow for Density Estimation

MAF: Inverse

- › Inverse mapping $x \mapsto z$:
 - › Compute all μ_i and α_i
 - › $z = \exp(-\alpha_1) \odot (x - \mu)$
- › Jacobian is lower diagonal, hence determinant can be computed efficiently.
- › Likelihood evaluation is easy and parallelizable.
- › Thus, the training is relatively fast.



MAF:results

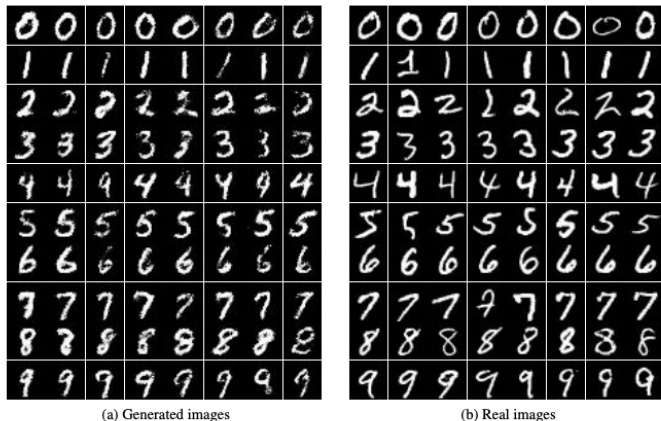
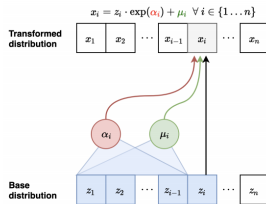


Figure 3: Class-conditional generated and real images from MNIST. Rows are different classes. Generated images are sorted by decreasing log likelihood from left to right.

Inverse Autoregressive Flow (IAF)

- › Forward mapping $z \mapsto x$:
 - › Sample all z_i ;
 - › Compute all μ_i and α_i .
 - › $x = \exp(-\alpha) \odot (z - \mu)$.
- › Inverse mapping $x \mapsto z$:
 - › Sequential calculation.
 - › $z_i = \exp(-\alpha_i(z_{<i}))(x - \mu_i(z_{<i}))$.
- › Fast to sample from, slow to evaluate likelihoods of data points (train).



IAF:results



(a) Generated images



(b) Real images

Figure 3: Class-conditional generated and real images from MNIST. Rows are different classes. Generated images are sorted by decreasing log likelihood from left to right.

MAF and IAF

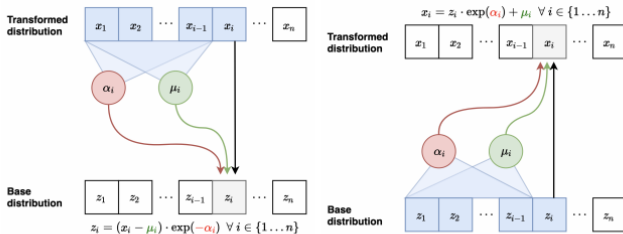


Figure: Inverse pass of MAF (**left**) vs. Forward pass of IAF (**right**)

- › MAF and IAF use autoregressive transformations based on MADE building block.
- › One can see that IAF forward mapping and MAF Inverse mapping are connected up to parameterisation.
- › In fact, they are inverse of each other.

MAF and IAF

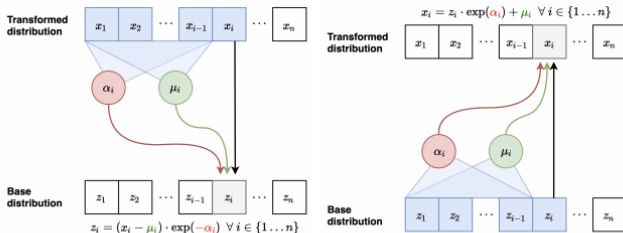
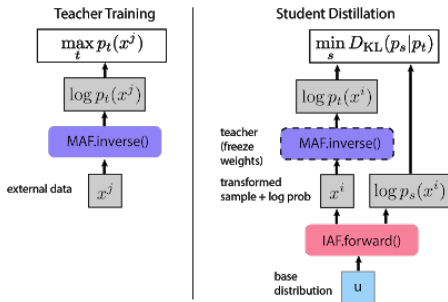


Figure: Inverse pass of MAF (**left**) vs. Forward pass of IAF (**right**)

- › MAF: Fast likelihood evaluation, slow sampling - best for training based on MLE, density estimation.
- › IAF: Fast sampling, slow likelihood evaluation - best for real-time generation.

Teacher-Student Model



- › Two part training with a teacher and student model.
- › Teacher (MAF) trained first, than student (IAF) initialised.
- › Student model cannot efficiently evaluate density for external datapoints but allows for efficient sampling.

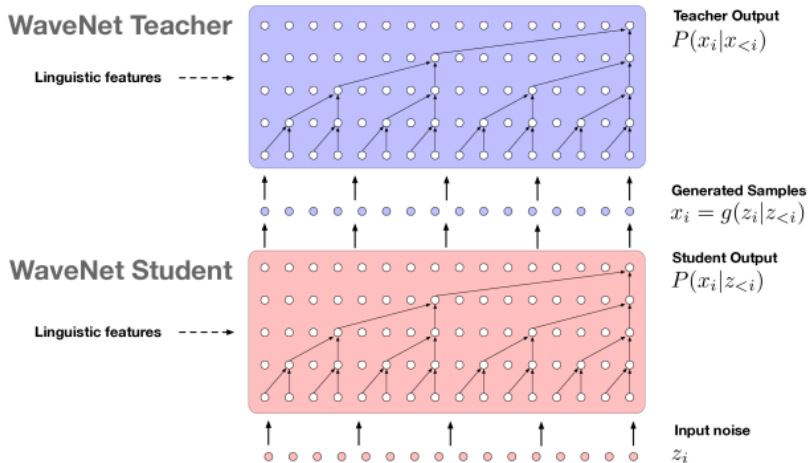
Probability Density Distillation

- › Student, s , is trained to match the teachers' distribution t using KL divergence:

$$KL(s, t) = \mathbb{E}_{x \sim s} [\log s(x) \log t(x)]$$

- › Training:
 - › Train teacher via MLE and obtain likelihood.
 - › Train student to minimize KL divergence.
 - › Use student to sample.
- › Improves sampling efficiencies by a factor 100 for Wavenet.

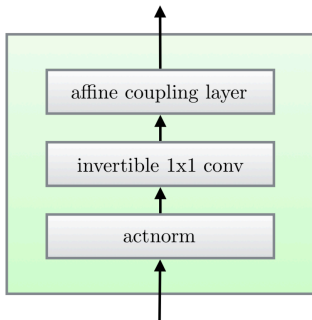
Parallel Wavenet



Gives fast and efficient in training algorithm for sound generation.

GLOW

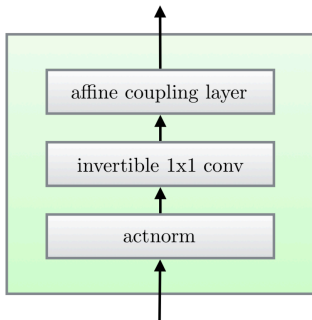
Generative Flow with Invertible 1x1 Convolutions



- › Updates NICE and RealNV and following their idea.
- › Uses block with several layers.

Blog by L. Weng
D. Kingma et al. Glow: Generative Flow with Invertible 1x1
Convolutions

GLOW: Layers

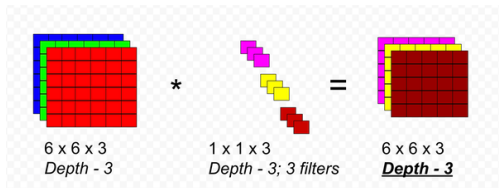


- › Activation normalization (short for “actnorm”): affine transformation using two trainable parameters (scale and bias).
- › Invertible 1x1 conv: generalization of any permutation (like r-NVP) of the channel ordering.
- › Affine coupling layer. Similar to rNVP.

Invertible 1x1 conv layer

- › We have an invertible 1x1 convolution:

$$f = \text{conv2D}(h, W).$$



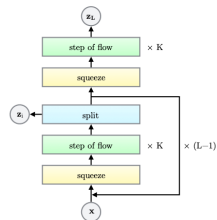
- › We need to compute the Jacobian determinant $|\det \partial f / \partial h|$.
- › In fact :

$$\log \left| \det \frac{\partial \text{conv2D}}{\partial h} \right| = \log(|\det W|^{h \cdot w}) = h \cdot w \log(|\det W|).$$

The latter operation can be computed using PU decomposition for matrix W of size $c \times c$ as $\mathcal{O}(c)$.

Summary of layers transformation

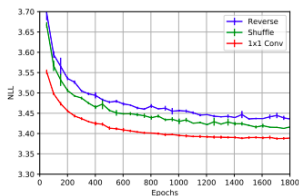
Description	Function	Reverse Function	Log-determinant
Actnorm. See Section 3.1.	$\forall i, j : \mathbf{y}_{i,j} = \mathbf{s} \odot \mathbf{x}_{i,j} + \mathbf{b}$	$\forall i, j : \mathbf{x}_{i,j} = (\mathbf{y}_{i,j} - \mathbf{b})/\mathbf{s}$	$h \cdot w \cdot \text{sum}(\log \mathbf{s})$
Invertible 1×1 convolution. $\mathbf{W} : [c \times c]$. See Section 3.2.	$\forall i, j : \mathbf{y}_{i,j} = \mathbf{W}\mathbf{x}_{i,j}$	$\forall i, j : \mathbf{x}_{i,j} = \mathbf{W}^{-1}\mathbf{y}_{i,j}$	$h \cdot w \cdot \log \det(\mathbf{W}) $ or $h \cdot w \cdot \text{sum}(\log \mathbf{s})$ (see eq. (10))
Affine coupling layer. See Section 3.3 and (Dinh et al., 2014)	$\mathbf{x}_a, \mathbf{x}_b = \text{split}(\mathbf{x})$ $(\log \mathbf{s}, \mathbf{t}) = \text{NN}(\mathbf{x}_b)$ $\mathbf{s} = \exp(\log \mathbf{s})$ $\mathbf{y}_a = \mathbf{s} \odot \mathbf{x}_a + \mathbf{t}$ $\mathbf{y}_b = \mathbf{x}_b$ $\mathbf{y} = \text{concat}(\mathbf{y}_a, \mathbf{y}_b)$	$\mathbf{y}_a, \mathbf{y}_b = \text{split}(\mathbf{y})$ $(\log \mathbf{s}, \mathbf{t}) = \text{NN}(\mathbf{y}_b)$ $\mathbf{s} = \exp(\log \mathbf{s})$ $\mathbf{x}_a = (\mathbf{y}_a - \mathbf{t})/\mathbf{s}$ $\mathbf{x}_b = \mathbf{y}_b$ $\mathbf{x} = \text{concat}(\mathbf{x}_a, \mathbf{x}_b)$	$\text{sum}(\log(\mathbf{s}))$



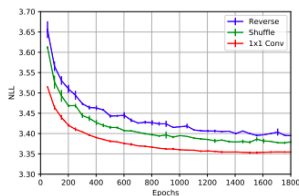
(b) Multi-scale architecture (Dinh et al., 2016).

Each layer is then followed by a subsequent multi-scale procedure (like it was done in NICE and r-NVP).

Additive vs Affine



(a) Additive coupling.



(b) Affine coupling.

Figure 3: Comparison of the three variants - a reversing operation as described in the RealNVP, a fixed random permutation, and our proposed invertible 1×1 convolution, with additive (left) versus affine (right) coupling layers. We plot the mean and standard deviation across three runs with different random seeds.

Authors claim that:

- › Affine faster than additive.
- › 1×1 convolution performs like better randomisation.

Unfortunately, to train on celeba, one needs a lot of GPU-days.

Sampling Temperature

- › In order to get more realistic sampling, one can use a reduced-temperature model.
- › In this work:

$$p_{\theta,T}(x) \sim p_{\theta}^{T^2}(x)$$

- › Temperature is a free parameter for sampling.

R. Dahl et al. Pixel Recursive Super Resolution

GLOW: Results



Figure 4: Random samples from the model, with temperature 0.7
Denis Derkach, Maksim Artemev, Artem Ryzhikov

Dependence on the Depth

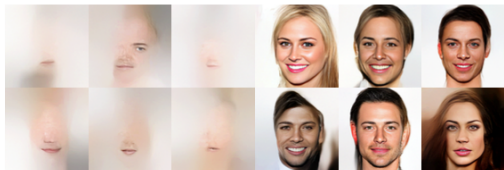


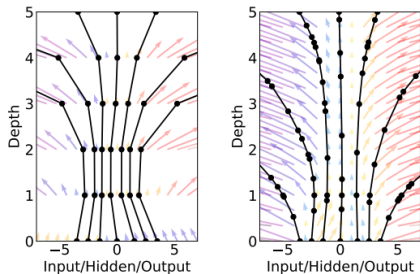
Figure 9: Samples from shallow model on left vs deep model on right. Shallow model has $L = 4$ levels, while deep model has $L = 6$ levels

Discussion

- › Adds additional parameters

FFJORD

Motivation



We can relax even more restrictions:

- › Do we really care of having discrete steps?
- › Can we change the Jacobian to something more stochastic?
- › We than thing of system of continuous-time dynamics.
- › This ideas led to a branch called NeuralODE.

Chen et al. Neural Ordinary Differential Equations

Continuous Normalizing Flows

Model the generative process with continuous dynamics:

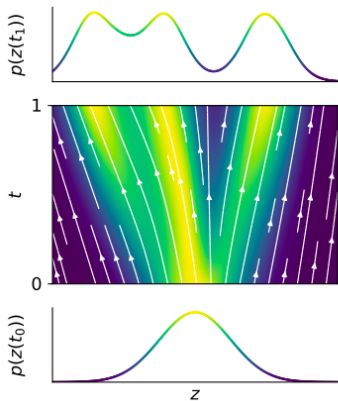
$$z_0 \sim p(z_0)$$

$$\frac{\partial z}{\partial t} = f_\theta(z_t, t)$$

$$x = z_t = z_0 + \int_{t_0}^{t_1} f_\theta(z_t, t) dt$$

To obtain the density we solve the initial value problem (IVP) under mild conditions:

$$\log(p_x) = \log(p_{z_0}) - \int_{t_0}^{t_1} \text{Tr} \frac{\partial f(z(t))}{\partial z(t)} dt$$



What this means

- › log-probability of the data under the discrete model:

$$\log p(x) = \log p(z_0) + \sum_{t=0}^T \log |\log \partial F^{-1} / \partial z_t|$$

- › log-probability of the data under continuous model:

$$\log(p_x) = \log(p_{z_0}) - \int_0^1 \text{Tr} \frac{\partial f(z(t))}{\partial z(t)} dt$$

- › sum of jacobian log-determinants \longrightarrow integral of jacobian trace.
This give $\mathcal{O}(N^3)$ caclulations.

Unbiased Log-Density Estimation

We can use stochastic trace estimation. For any matrix A and a distribution $p(e)$ over vectors where $\mathbb{E}[e] = 0$, $\text{cov}[e] = I$, we used Hutchinson's estimator:

$$\text{Tr}(A) = \mathbb{E}_{p(e)}[e^T A e]$$

Which brings us to calculable:

$$\log(p_x) = \log(p_{z_0}) - \mathbb{E}_{p(e)} \int_{t_0}^{t_1} e^T \frac{\partial f(z(t))}{\partial z(t)e} dt$$

The existence and uniqueness of solution requires that f and its first derivatives be Lipschitz continuous and can be calculated in $\mathcal{O}(N)$.

Training with adjoint Backprop

- › We need to compute $\partial L / \partial \theta$.
- › Given scalar objective:

$$\mathcal{L}(z_1) = \mathcal{L} \left(\int_0^1 f(z(t), t, \theta) dt \right)$$

- › we can obtain $\partial \mathcal{L} / \partial \theta$ for gradient-based optimization by solving another IVP.

We define a new quantity, the adjoint, a_t , which has dynamics $\frac{\partial a_t}{\partial t}$

$$a_t = -\frac{\partial L}{\partial z_t} \quad \frac{\partial a_t}{\partial t} = -a_t^T \frac{\partial f(z_t, t, \theta)}{\partial z_t}$$

then solving backwards in time gives the desired gradients of the loss with respect to the parameters

$$\frac{\partial L}{\partial \theta} = \int_{t_1}^{t_0} a_t^T \frac{\partial f(z_t, t, \theta)}{\partial \theta} dt$$

This allows us to use a black-box ODE solver to compute z_1 and also $\partial L / \partial \theta$.

FFJORD: results

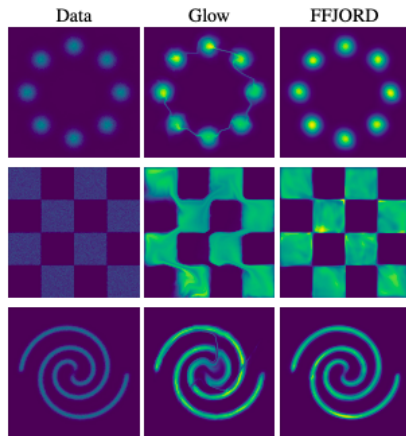


Figure 2: Comparison of trained FFJORD and Glow models on 2-dimensional distributions including multi-modal and discontinuous densities.

FFJORD: discussion

› Advantages

- › Guaranteed inverse by reversing order of integration, regardless of model parameterization
- › Efficient, unbiased log-probability estimation without restricting the Jacobian of the transformation
- › Does not require dimension splitting or ordering choices
- › Reversible generative models can now be defined with standard neural network architectures

› Disadvantages

- › Must rely on adaptive numerical ODE solvers for stable training
- › Computation time determined by solver, not user
- › Currently 4-5x slower than other reversible generative models (Glow, Real-NVP)

Conclusion

Method		Train on data	One-pass Sampling	Exact log-likelihood	Free-form Jacobian
Change of Variables	Variational Autoencoders	✓	✓	✗	✓
	Generative Adversarial Nets	✓	✓	✗	✓
	Likelihood-based Autoregressive	✓	✗	✓	✗
	Normalizing Flows	✗	✓	✓	✗
	Reverse-NF, MAF, TAN	✓	✗	✓	✗
	NICE, Real NVP, Glow, Planar CNF	✓	✓	✓	✗
	FFJORD	✓	✓	✓	✓

Table 1: A comparison of the abilities of generative modeling approaches.